

FAB - RAG Chatbot

1. Project Overview

FAB-RAG is a lightweight and efficient Retrieval-Augmented Generation (RAG) system specifically engineered for interactive question answering over a custom dataset consisting of textual content about the history of Ancient Greece. Leveraging the capabilities of LangChain for orchestration, FAISS for fast similarity search, and OpenAI's `gpt-4o-mini` for natural language generation, FAB-RAG delivers an intuitive and responsive user experience. One of its key strengths lies in maintaining conversation history through session memory, allowing it to answer follow-up questions contextually, thereby simulating human-like dialogue.

2. Objective

The primary goal of FAB-RAG is to build a robust, conversationally aware QA system with the following objectives:

- **Answer Accuracy:** Extract and deliver precise answers from curated historical text files.
- **Conversational Context Retention:** Maintain session history to handle follow-up or related queries effectively.
- **Faithful Contextual Responses:** Ensure answers are grounded strictly in the retrieved documents to avoid hallucinations.

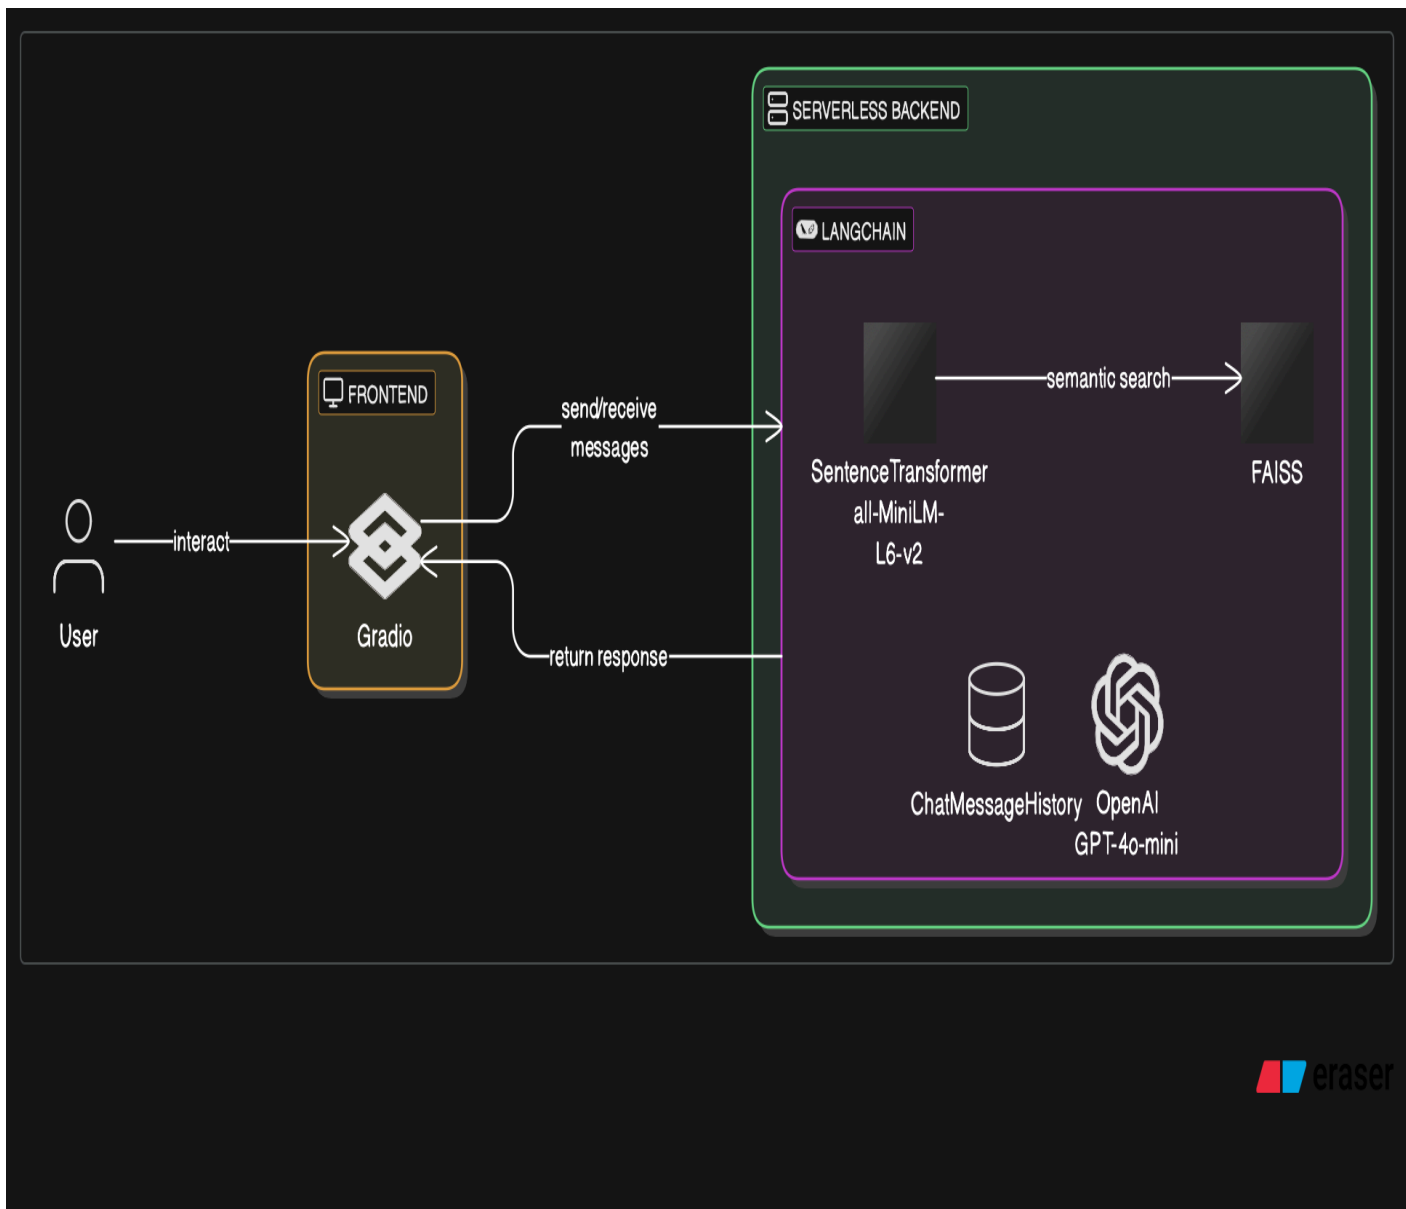
3. Technology Stack

The solution is developed using a highly modular and flexible stack:

- **Framework:** Python 3.12.9 and LangChain
 - Python provides the language runtime, while LangChain offers essential abstractions to create LLM-based applications with chaining, retrieval, and memory components.
- **LLM (Language Model):** OpenAI `gpt-4o-mini`
 - This is a lightweight variant of OpenAI's GPT-4 model, optimized for low latency and cost-effective usage without compromising significantly on reasoning quality.
- **Embeddings:** SentenceTransformer `all-MiniLM-L6-v2`
 - A 384-dimensional model from Hugging Face that offers high-quality sentence-level embeddings ideal for semantic similarity tasks. This model strikes a balance between accuracy, speed, and compute efficiency.
- **Vector Store:** FAISS (Facebook AI Similarity Search)

- FAISS enables rapid approximate nearest-neighbor search on dense vectors, which is crucial for scalable and performant document retrieval.
- **Frontend:** Gradio
 - Provides an easy-to-use and responsive web interface for users to ask questions and view responses in real-time.
- **Session Store:** LangChain ChatMessageHistory
 - Manages individual user sessions, preserving conversational history that enhances context understanding in subsequent interactions.

4. Architecture



5. Key Modules

5.1 Document Loader

- Loads `.txt` files from the `ancient_greece` folder.
- Cleans and parses files into LangChain Document objects.
- **Note:** Text splitter is not used currently. Each `.txt` file is treated as a single document, which may affect retrieval granularity.
- **Ideal Strategy (Future Plan):** Integrate a recursive text splitter with chunking and overlap to improve retrieval precision and manage long documents efficiently.

5.2 FAISS Indexing

- Builds or loads FAISS index.
- Uses cosine similarity for vector retrieval.

5.3 Embedding Model: `all-MiniLM-L6-v2`

- 384-dimensional transformer model.
- Chosen for its balance of performance, speed, and dimensionality.
- Trained on NLI, STS, and paraphrasing tasks — ideal for semantic search.

5.4 LLM Model: `gpt-4o-mini`

- Offers concise, fluent generation.
- Handles historical context, reasoning, and rewriting of follow-up queries.

5.5 Contextual Retriever

- Reformulates follow-up questions using a history-aware chain.
- Ensures RAG retrieval context is aligned with the actual intent.

5.6 Conversational Memory

- Maintains separate sessions using unique session IDs.
- Built using `RunnableWithMessageHistory`.

5.7 Evaluation

- Metrics computed: Cosine Similarity, BLEU, ROUGE-1/2/L, METEOR, BERTScore.
- LLM-evaluated: Faithfulness, Relevance, Context-Faithfulness, Truthfulness.

5.8 Customer Satisfaction Evaluation

To assess real-world effectiveness, customer satisfaction should be evaluated through user feedback collected during A/B Testing & afterwards for further improvement. Users can be asked to rate their experience based on the following metrics:

- **Accuracy of Answers:** Did the chatbot provide correct and helpful responses?
- **Ease of Use:** Was the interface intuitive and easy to interact with?
- **Context Awareness:** Did the chatbot understand follow-up questions and maintain conversation flow?
- **Overall Satisfaction:** General impression and likelihood to recommend the system to others.

6. Evaluation Snapshot

Sample QA :

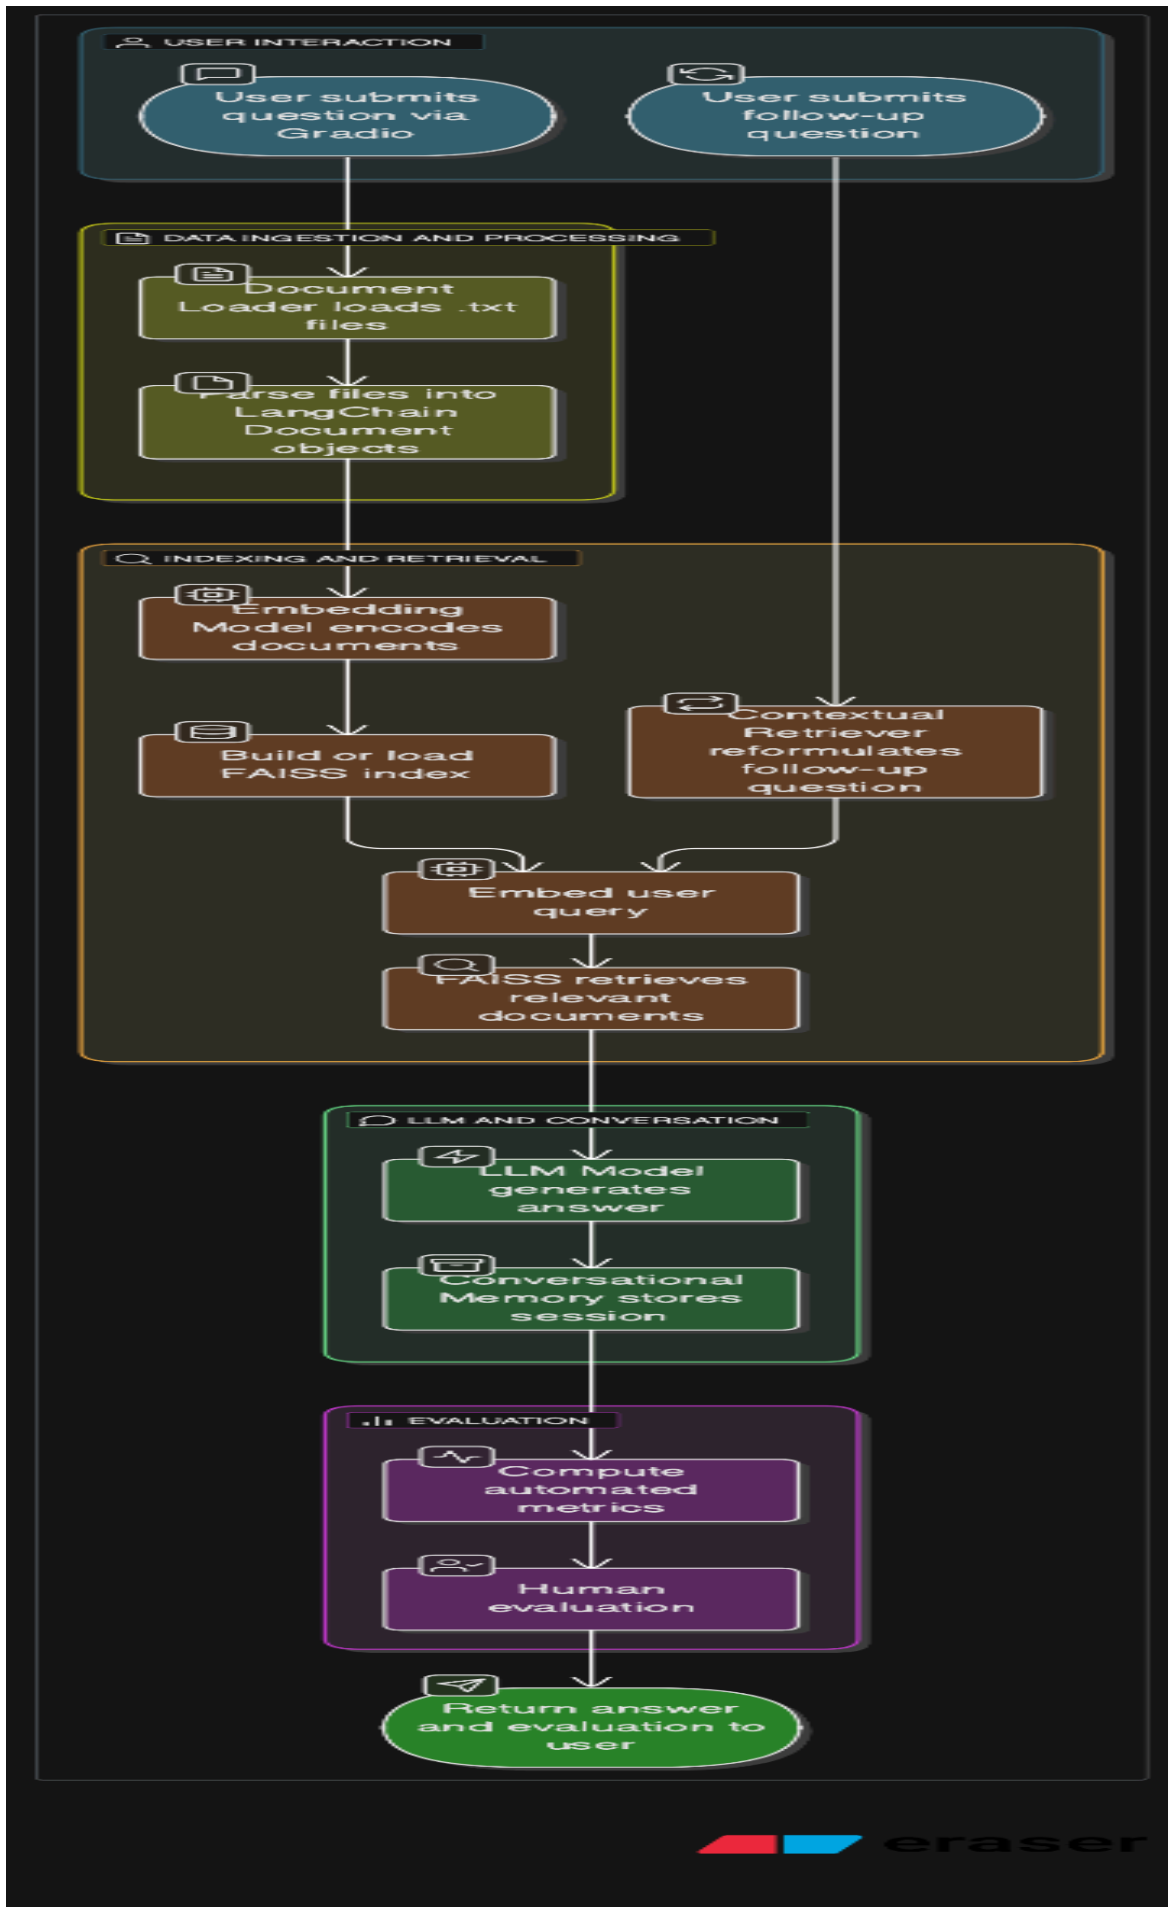
```
[
  {
    "question": "When did the story of Ancient Greece begin?",
    "expected_answer": "The story begins around 3000 BCE when a group of seafaring people known as the Minoans settled on the island of Crete.",
    "ground_truth_chunk": "The story begins around 3000 BCE when a group of seafaring people known as the Minoans settled on the island of Crete."
  }
]
```

Sample Result :

question,expected_answer,ground_truth_chunk,rag_answer,cosine_similarity,bleu_score,rouge-1,rouge-2,rouge-l,meteor,bertscore_f1,faithfulness,context_faithfulness,relevance,truthfulness

How does the internal Pipeline will look like

Below the the internal architecture of pipelines



7. Reasoning for Model Choices

Embedding Model: **all-MiniLM-L6-v2**

- **Dimensionality:** 384
- **Reasoning:**
 - Lightweight and fast inference, suitable for CPU and local deployment.
 - Strong semantic representation trained on natural language inference (NLI), semantic textual similarity (STS), and paraphrasing tasks, making it ideal for semantic search and retrieval.
 - Provides a balanced embedding dimension (384) that offers good accuracy without excessive computational cost.
 - Well-integrated with FAISS and LangChain, ensuring seamless vector store creation and retrieval.

Discussion on Embedding Alternatives

Why Not **text-ada-002** Embeddings?

- **text-ada-002** embeddings, offered by OpenAI, have higher dimensionality and tend to be more computationally expensive and costly.
- While powerful, they are often overkill for smaller or domain-specific datasets like Ancient Greece historical texts.
- **all-MiniLM-L6-v2** provides sufficiently rich semantic information at lower latency and cost.
- Additionally, SentenceTransformer models like **all-MiniLM-L6-v2** can be used offline, offering flexibility not available with API-based embeddings like **text-ada-002**.

Why Not BERTopic Embeddings?

- BERTopic embeddings are primarily designed for topic modeling and clustering rather than fine-grained semantic similarity search.
- The clustering mechanism in **BERTopic** can reduce retrieval precision in tasks that require detailed passage-level matching, such as question answering.
- The added computational complexity of BERTopic is unnecessary for a RAG system focused on accurate document retrieval.
- SentenceTransformer embeddings, such as **all-MiniLM-L6-v2**, better align with FAISS vector stores for efficient and accurate similarity-based retrieval.

LLM: **gpt-4o-mini**

- Offers cost-effective, fluent, and context-aware generation.
- Handles conversational history and complex historical context effectively.
- Balances inference speed with generation quality for an interactive QA system.

8. Deployment & Usage

Run locally using **gradio** or integrate with web frontends.

Supports:

- Dynamic question input
- Session-based memory
- Document-based QA with RAG

Want to check out Code and Deployment, Check it out below:

- HuggingFace : https://huggingface.co/spaces/Nishthaaa/langchain_openai_rag_chatbot
- Github : <https://github.com/kukretinishtha/fab-rag>

9. Future Work

- Add PDF and image OCR ingestion
- Integrate multilingual support
- Expand domain to other historical datasets
- Add text splitting for more granular retrieval using recursive chunking with overlap
- Extensive evaluation and report need to be generated as this project demonstrates how the chatbot should be in a production ready environment.

10. Contributors

- [@kukretinishtha](#)